

ENHANCING **METRIC PRIVACY** WITH A **SHUFFLER**

Andreas Athanasiou

Catuscia Palamidessi, Kostas Chatzikokolakis



“Have you cheated on the exam?”

—A random (angry) professor

Have you cheated on the exam?

Answers:
(using Differential Privacy)



Class A



Class A without Bob



20%
Yes



~20%
Yes

Bob's answer does not (really) change the result!

Differential Privacy (ϵ, δ)

	Same	Different	
Dataset X :	1, 15, 2, 3 , 50 , ... , 5		
Dataset X' :	1, 15, 2, 3 , 50 , ... , 90		

Adjacent Datasets

$$\mathbb{P}[M(X) \in S] \leq e^\epsilon \cdot \mathbb{P}[M(X') \in S] + \delta$$

ϵ = Privacy Loss

δ = (Rare) cases where
 ϵ doesn't hold

The probability to see the same outcome between every possible set of adjacent datasets:

- Is at most e^ϵ
- Can be above e^ϵ in at most δ % of the cases (for $0 \leq \delta \leq 1$)

Differential Privacy (ϵ, δ)

	Same	Different	
Dataset X:	1, 15, 2, 3 , 50 , ... , 5		
Dataset X':	1, 15, 2, 3 , 50 , ... , 90		

Adjacent Datasets

$$\mathbb{P}[M(X) \in S] \leq e^\epsilon \cdot \mathbb{P}[M(X') \in S] + \delta$$

ϵ = Privacy Loss

δ = (Rare) cases where ϵ doesn't hold

The probability to see the same outcome between every possible set of adjacent datasets:

- Is at most e^ϵ
- Can be above e^ϵ in at most δ % of the cases (for $0 \leq \delta \leq 1$)

Example:

$\epsilon = 0.2$ and $\delta = 0.01$:

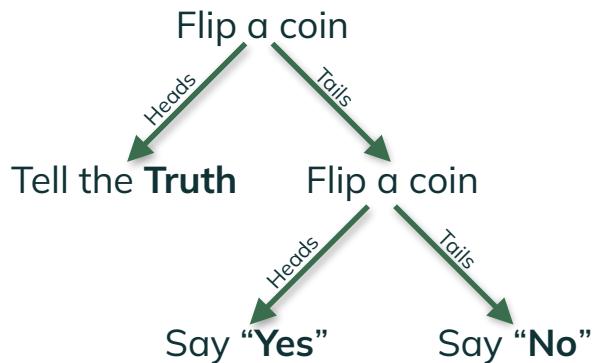
the mechanism offers 0.2 privacy, in 99% of cases



Randomised Response

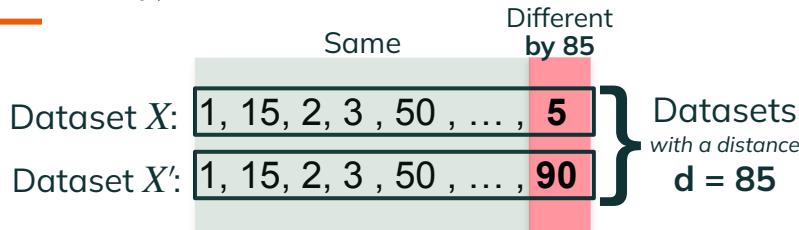
Question: “Have you cheated on the exam?”

Answer:



Metric Privacy (ϵ, δ)

(aka d-privacy)



$$\forall x \in X, x' \in X': \mathbb{P}[M(x) \in S] \leq e^{\epsilon \cdot d_x(x, x')} \mathbb{P}[M(x') \in S] + \delta$$

$\epsilon \cdot d$ = Privacy Loss

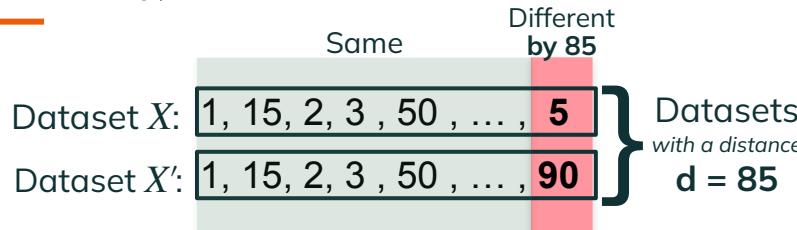
δ = (Rare) cases where $\epsilon \cdot d$ doesn't hold

The probability to see the same outcome between every possible set of datasets **that have a distance d** :

- Is at most $e^{\epsilon \cdot d}$
- Can be above $e^{\epsilon \cdot d}$ in at most δ % of the cases (for $0 \leq \delta \leq 1$)

Metric Privacy (ϵ, δ)

(aka d-privacy)



$$\forall x \in X, x' \in X': \mathbb{P}[M(x) \in S] \leq e^{\epsilon \cdot d_x(x, x')} \mathbb{P}[M(x') \in S] + \delta$$

$\epsilon \cdot d$ = Privacy Loss

δ = (Rare) cases where $\epsilon \cdot d$ doesn't hold

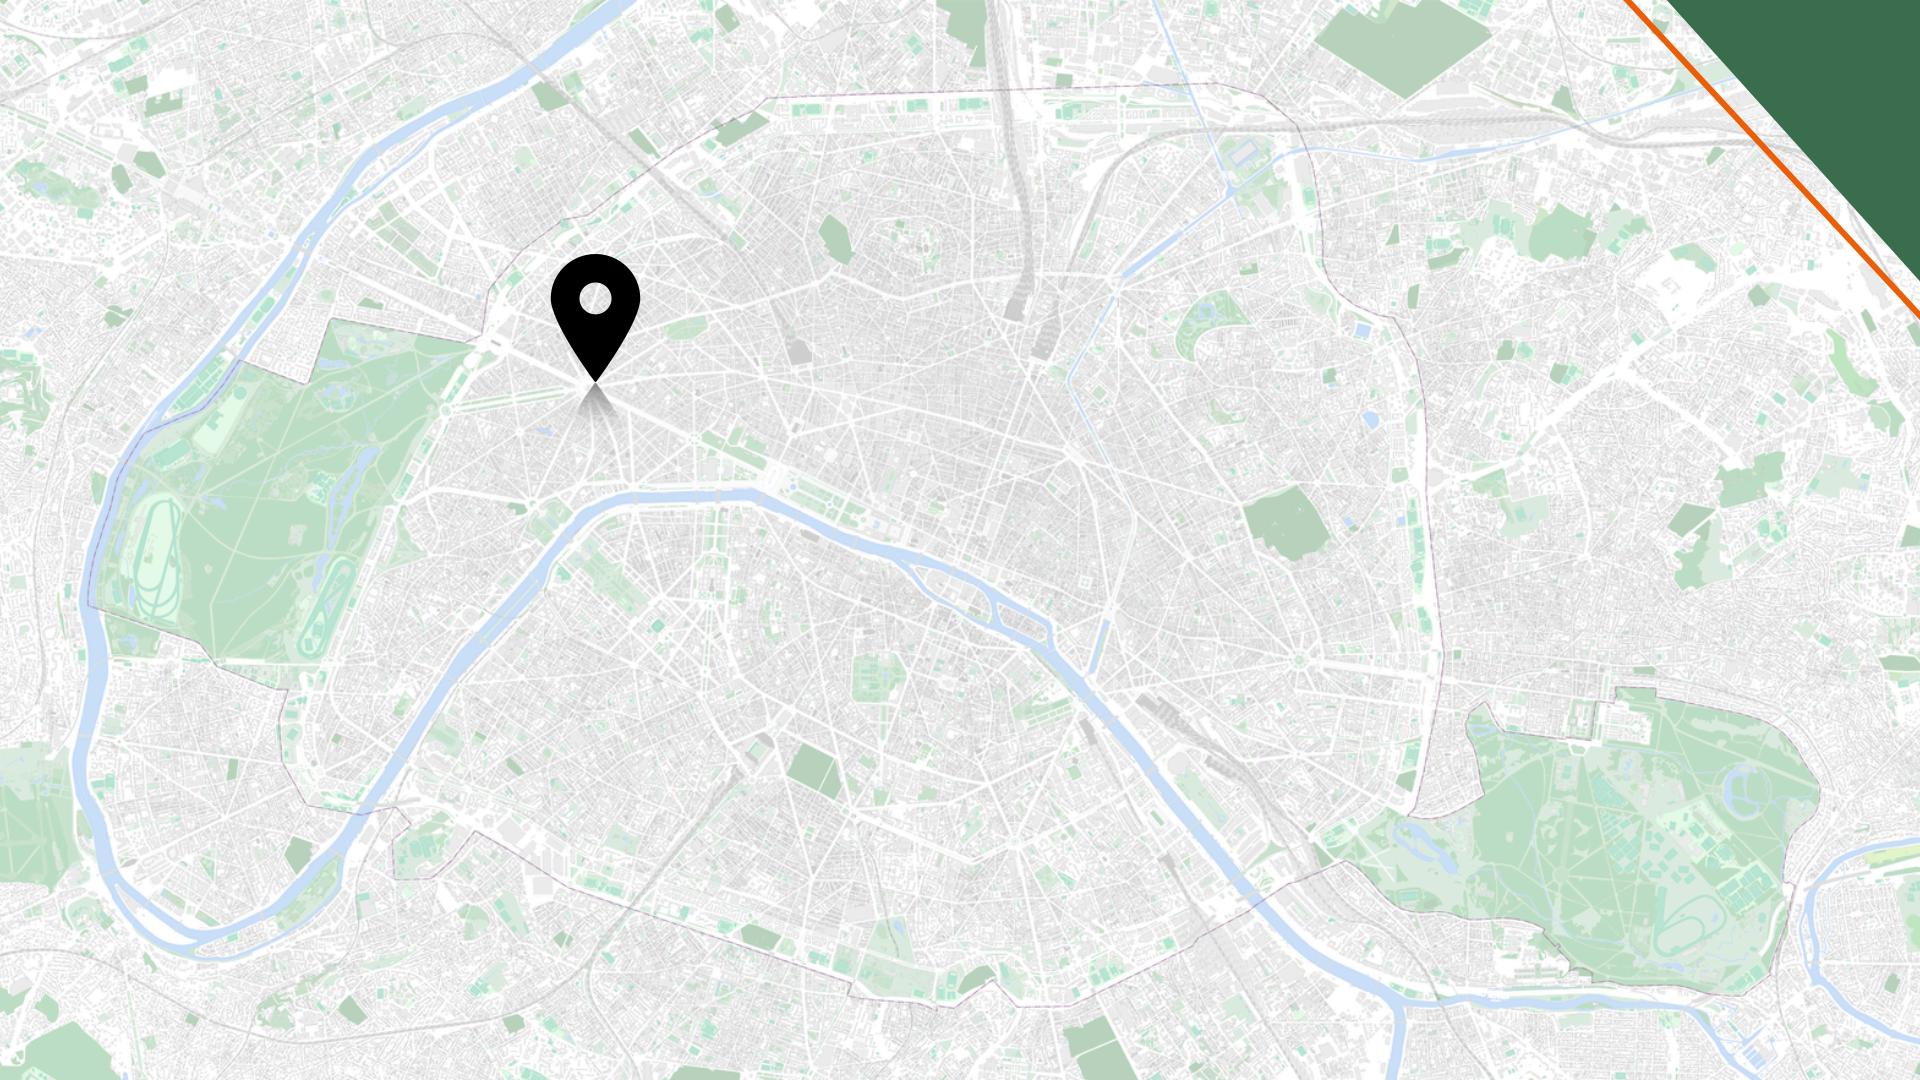
The probability to see the same outcome between every possible set of datasets **that have a distance d** :

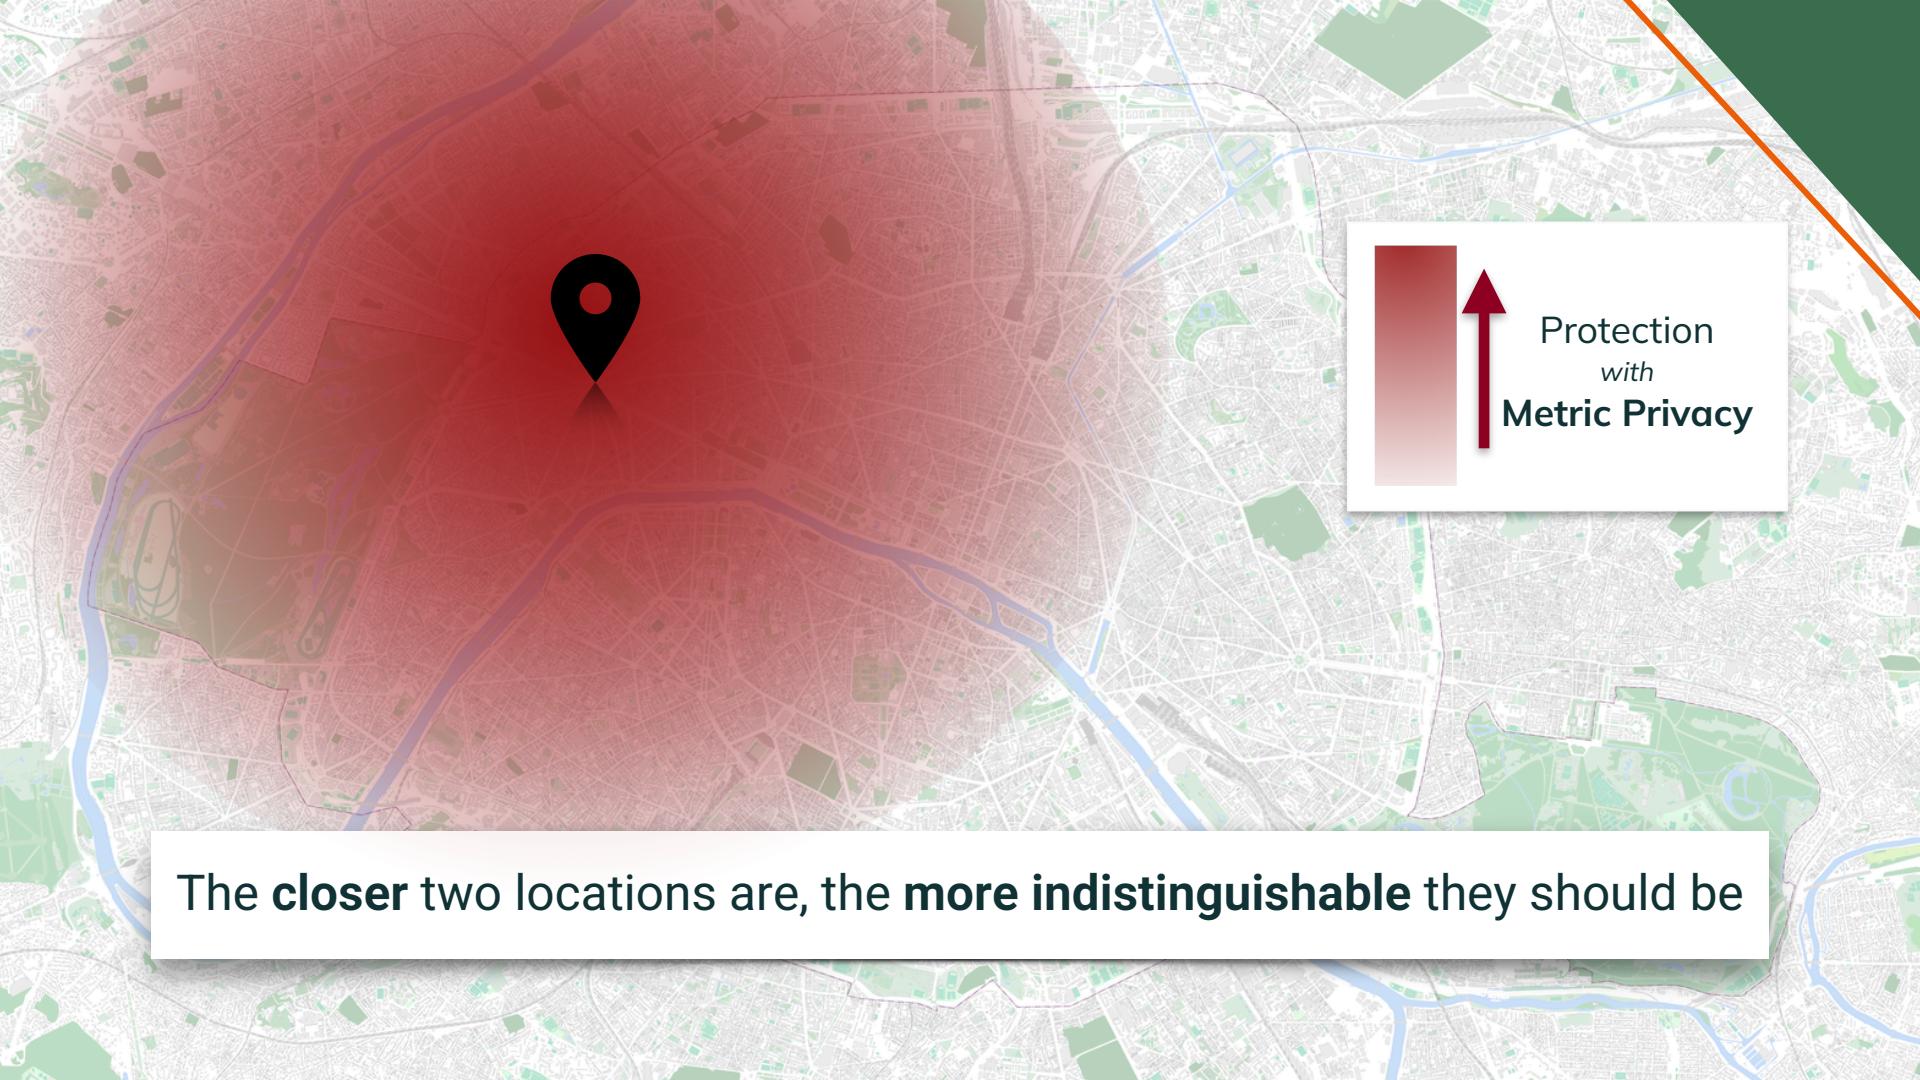
- Is at most $e^{\epsilon \cdot d}$
- Can be above $e^{\epsilon \cdot d}$ in at most δ % of the cases (for $0 \leq \delta \leq 1$)

Example:

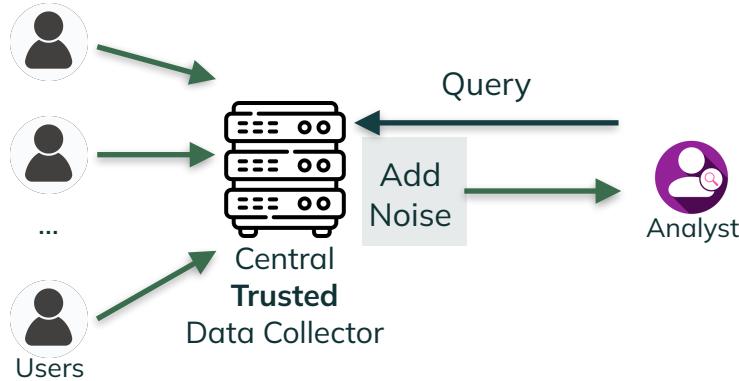
if $\epsilon = 0.2$ and $\delta = 0.01$:

the mechanism offers $0.2 \cdot d$ metric privacy (for every dataset with a distance d), in 99% of cases



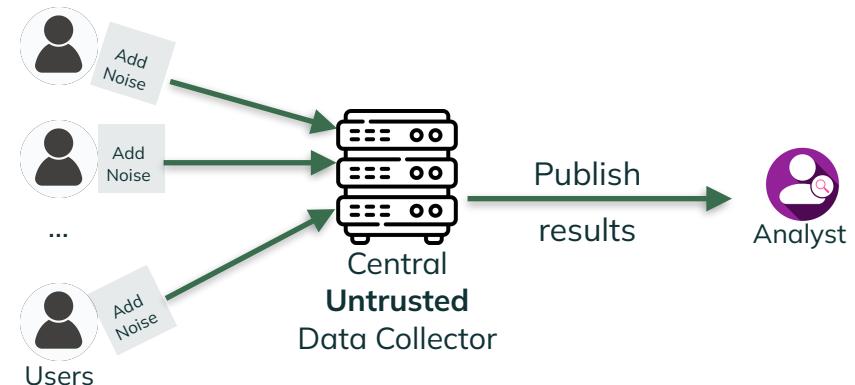
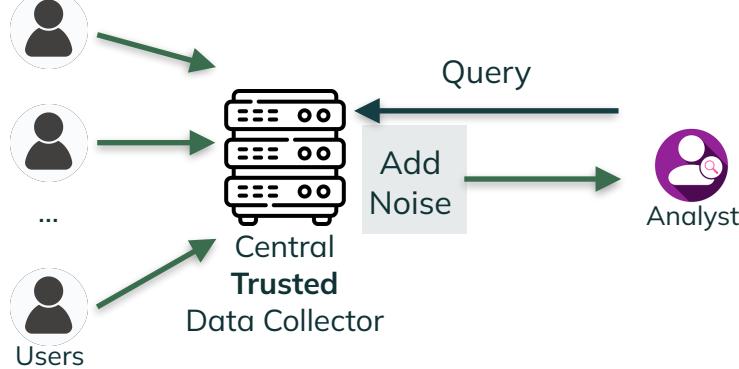


Models of Privacy

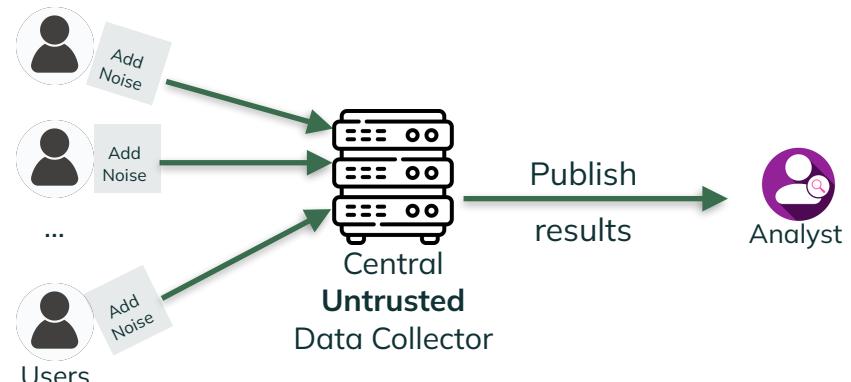
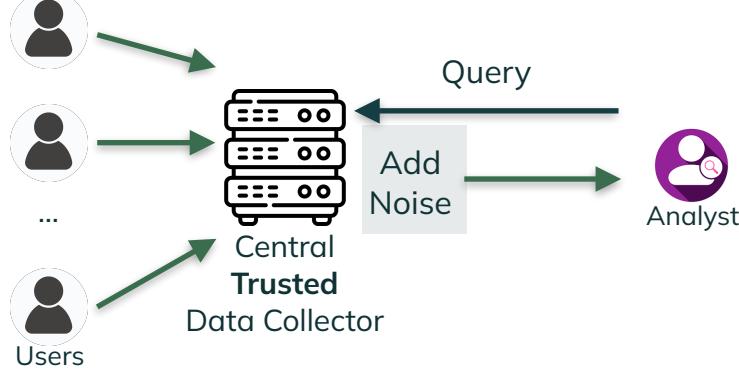


Central Model

Models of Privacy



Models of Privacy



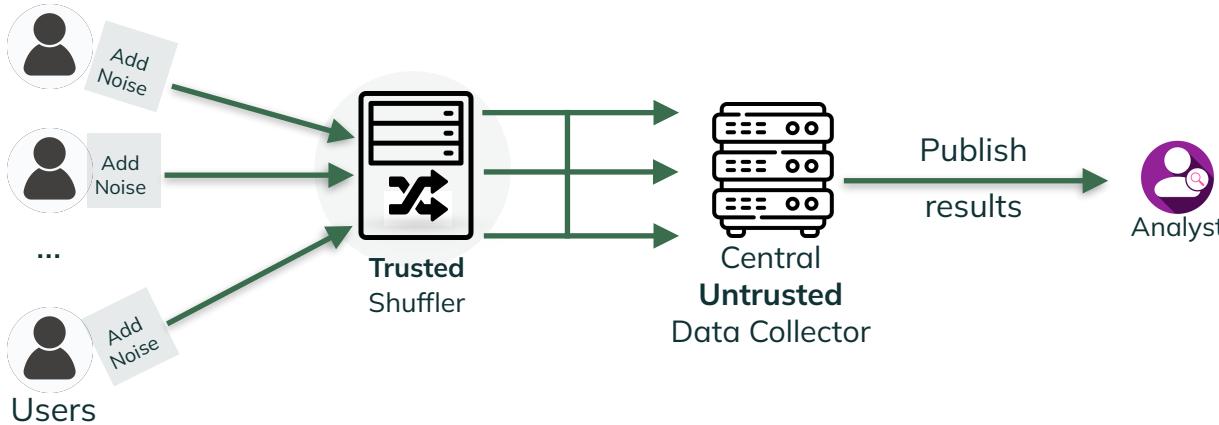
Central Model

- + Better utility
- Must trust the data collector

Local Model

- + No need to trust a central entity
- Worse utility

The Shuffle Model



Shuffle Model

- + Better utility than the Local Model
- + Less trust than the Central Model

Metric Privacy *in the* Shuffle Model

Problem: Private summation of integers
 n users | each user i has a value $x_i \in \{0, 1, \dots, k\}$ for $k \in \mathbb{N}$

Metric Privacy *in the* Shuffle Model

Problem: Private summation of integers
 n users | each user i has a value $x_i \in \{0, 1, \dots, k\}$ for $k \in \mathbb{N}$

Contributions:

-  - Shuffle

Randomised Response mechanism

-  - Shuffle

Geometric mechanism

-  - Shuffle

Symmetric Generalised Discrete Laplace distribution

RR - Shuffle

n users | each user i has a value $x_i \in \{0, 1, \dots, k\}$



x_1



x_2

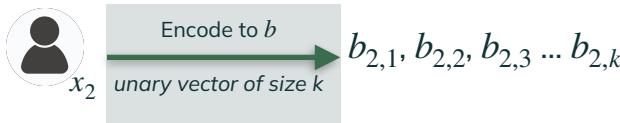
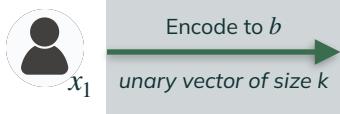
...



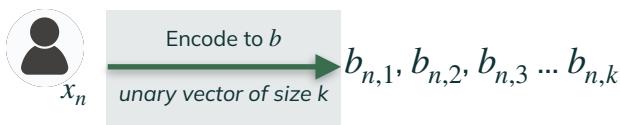
x_n

RR - Shuffle

n users | each user i has a value $x_i \in \{0, 1, \dots, k\}$



...



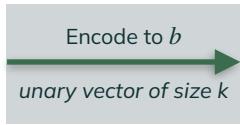
$$\sum_i b_i = x_i$$

RR - Shuffle

n users | each user i has a value $x_i \in \{0, 1, \dots, k\}$



x_1



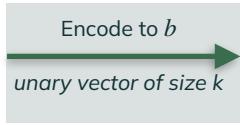
Run



$\forall b_i$



x_2



Run



$\forall b_i$

...

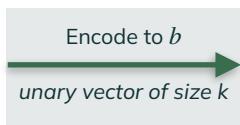
...

...

...



x_n



Run

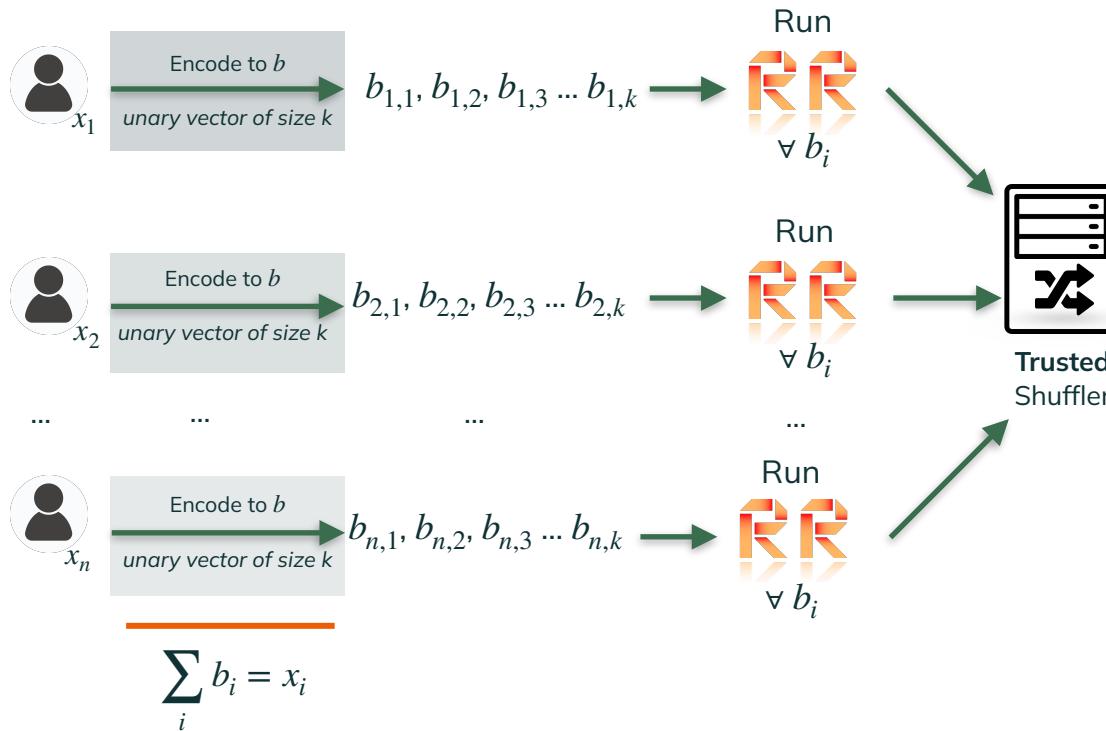


$\forall b_i$

$$\sum_i b_i = x_i$$

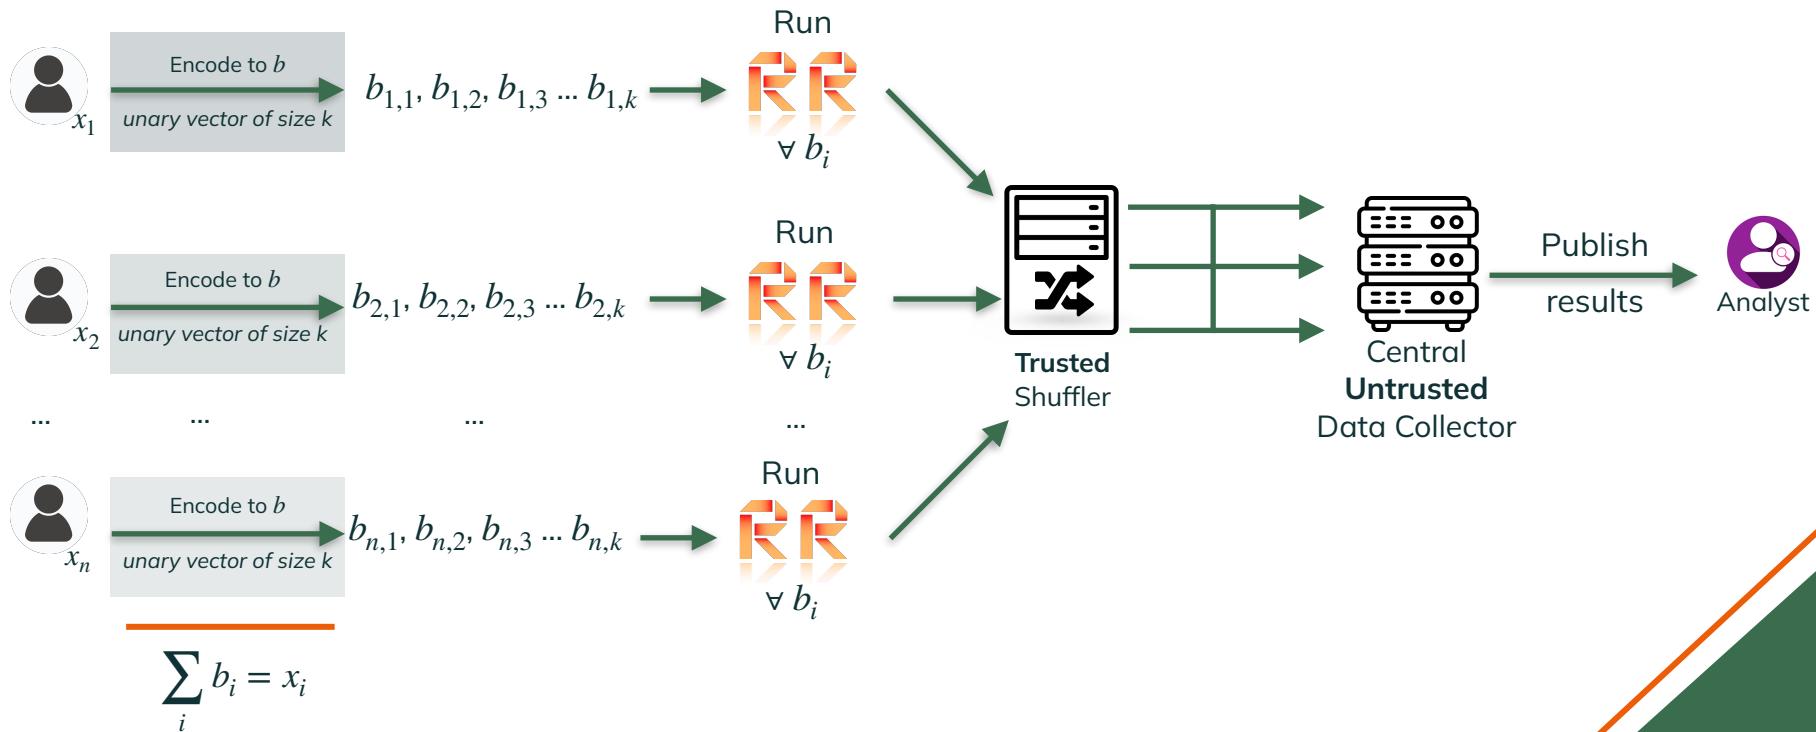
RR - Shuffle

n users | each user i has a value $x_i \in \{0, 1, \dots, k\}$



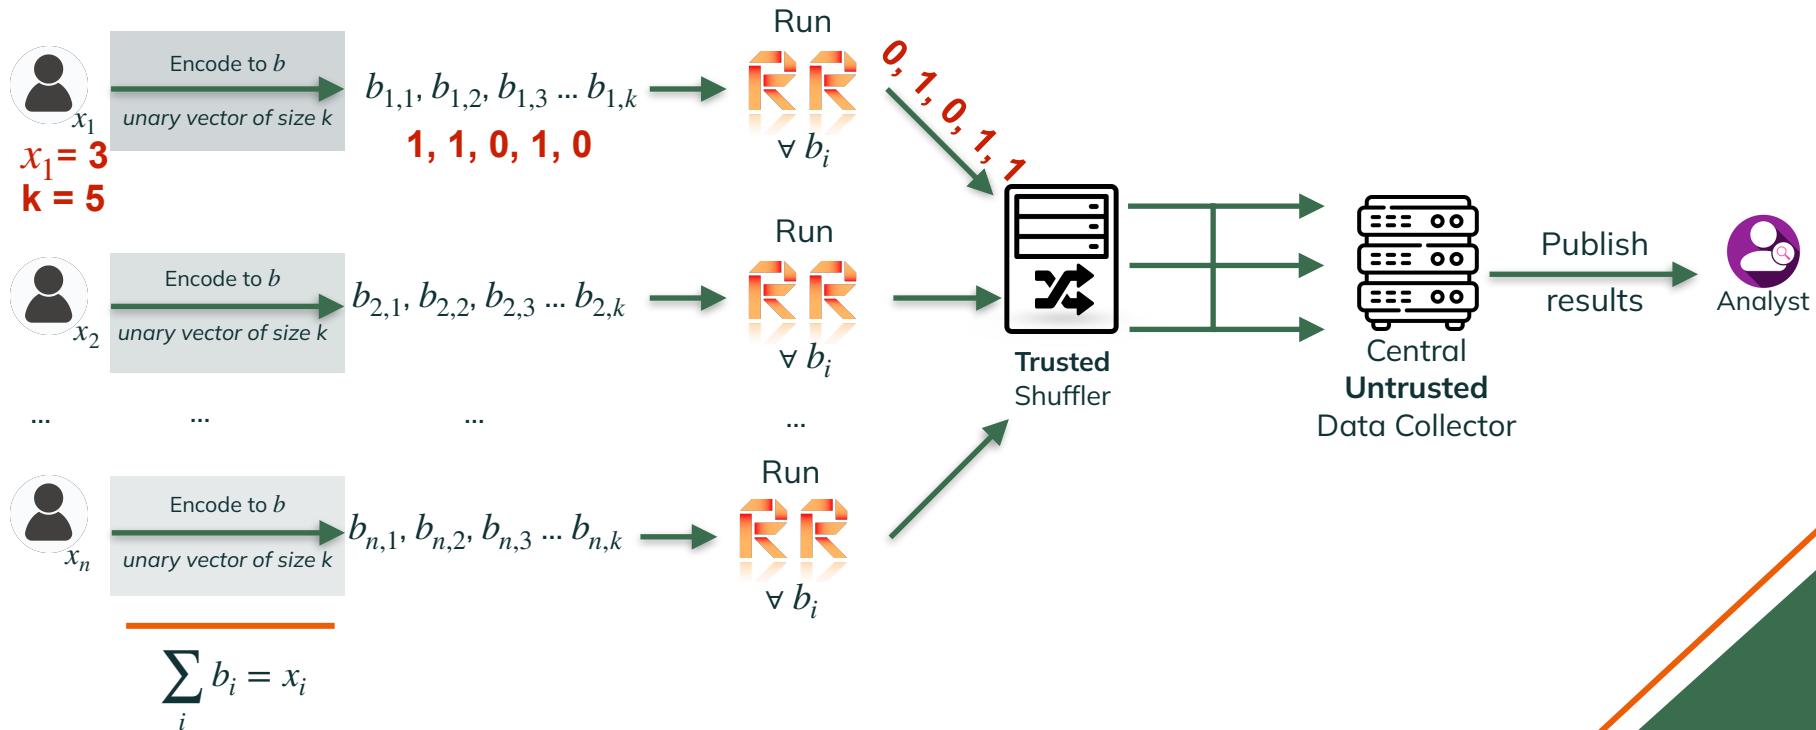
RR - Shuffle

n users | each user i has a value $x_i \in \{0, 1, \dots, k\}$



RR - Shuffle

n users | each user i has a value $x_i \in \{0, 1, \dots, k\}$



RR - Shuffle: Privacy



Shuffle Model Property

Shuffling unary bits

is privacy-wise equivalent to

Summing unary bits

RR - Shuffle: Privacy

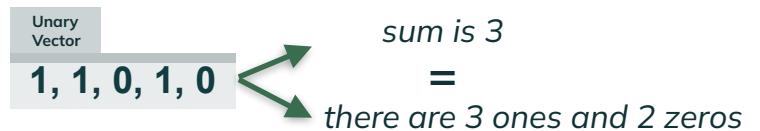


Shuffle Model Property

Shuffling unary bits

is privacy-wise equivalent to

Summing unary bits



RR - Shuffle: Privacy



Shuffle Model Property

Shuffling unary bits
is privacy-wise equivalent to
Summing unary bits



Binomial Distribution

The sum of all the random bits follows the Binomial Distribution.

RR-Shuffle needs a minimum number of users!

RR - Shuffle: Privacy



Shuffle Model Property

Shuffling unary bits
is privacy-wise equivalent to
Summing unary bits



Binomial Distribution

The sum of all the random bits follows the Binomial Distribution.

RR-Shuffle needs a minimum number of users!



Rare cases

- The **number of random bits** is *unexpectedly “small/large”*
- The **result of random bits** is *unexpectedly “small/large”*

Other Mechanisms

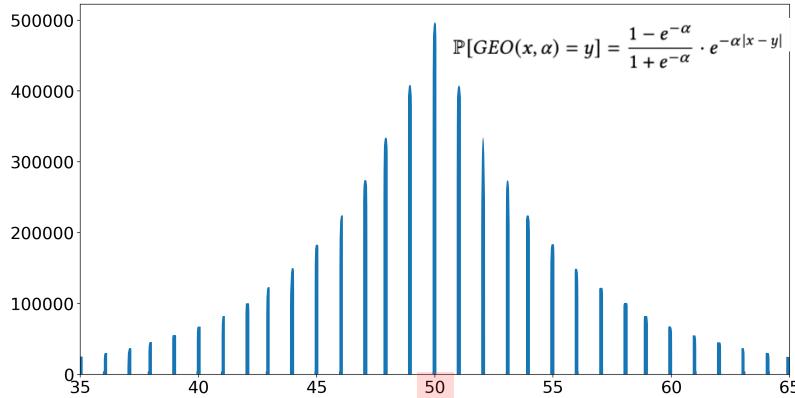
geo - Shuffle

Users sample noise from:

Geometric Mechanism

Applied with a parameter α to a user's value x

Produces y with exponentially decreasing probability wrt $d(x, y)$



Histogram:

Reported values of the Geometric Mechanism with parameter $\alpha = 0.2$ and input value $x = 50$

Other Mechanisms

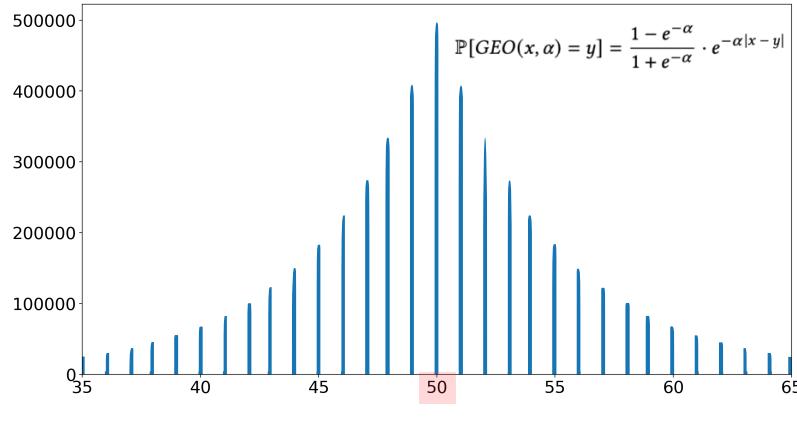
geo - Shuffle

Users sample noise from:

Geometric Mechanism

Applied with a parameter α to a user's value x

Produces y with exponentially decreasing probability wrt $d(x, y)$



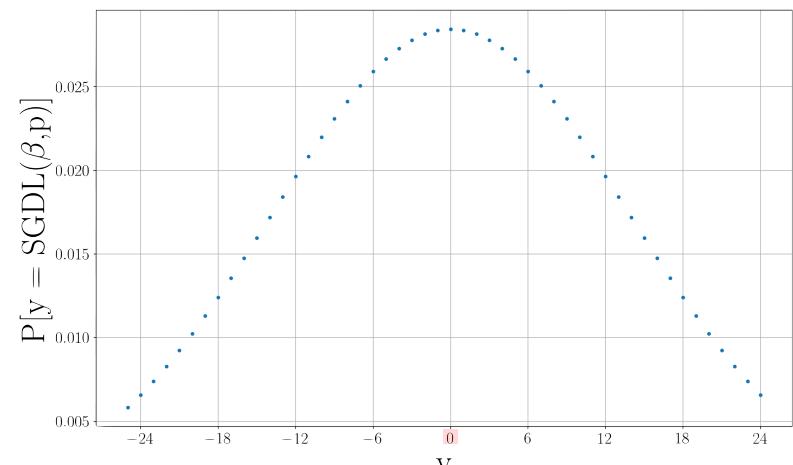
Histogram:
Reported values of the Geometric Mechanism with parameter $\alpha = 0.2$ and input value $x = 50$

SGDL - Shuffle

Users sample noise from:

Symmetric Generalised Discrete Laplace distribution (β, p)

:= difference between two **Negative Binomial** distributions ($\beta, 1 - p$)



PMF of SGDL:
Probability to report y for $\beta = 50, p = 0.5$

Other Mechanisms

geo - Shuffle

Users sample noise from:

Geometric Mechanism



Unary Encode



Shuffle

SGDL - Shuffle

Users sample noise from:

Symmetric Generalised Discrete Laplace distribution (β, p)



Unary Encode



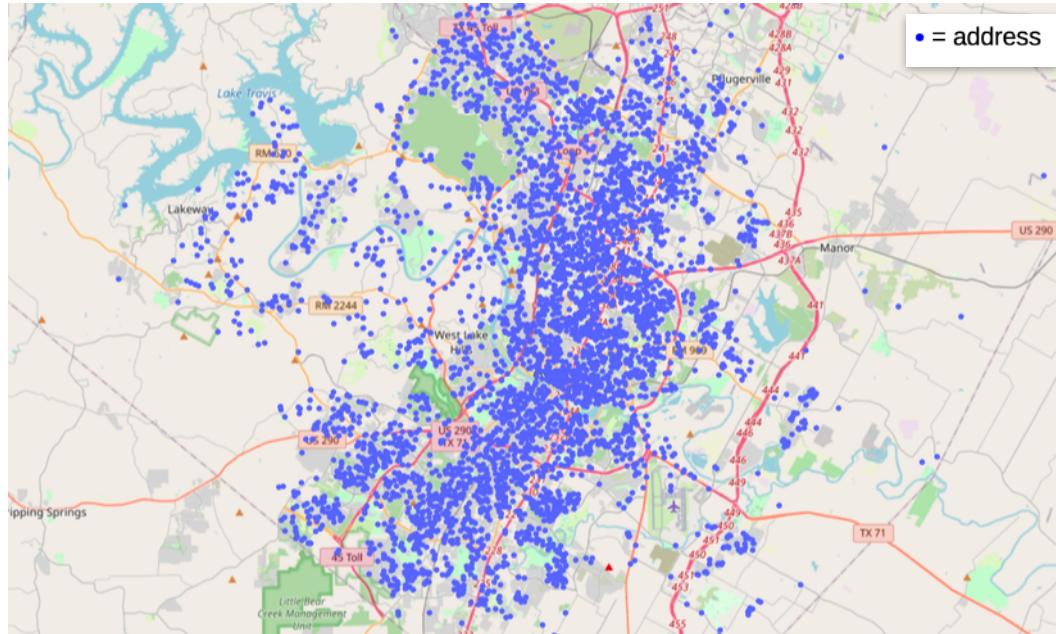
Shuffle

Comparison of Mechanisms

 - Shuffle	<ul style="list-style-type: none">Simple	<ul style="list-style-type: none">“Not great, not terrible” utilityNeeds a minimum number of users
geo - Shuffle	<ul style="list-style-type: none">Excellent utilityMedium trust on the shuffler: <i>the protocol retains some privacy even if the shuffler has been compromised</i>	<ul style="list-style-type: none">Not optimal utility
SGDL - Shuffle	<ul style="list-style-type: none">Optimal utility	<ul style="list-style-type: none">Heavy trust on the shuffler: <i>the protocol provides almost no privacy if the shuffler has been compromised</i>

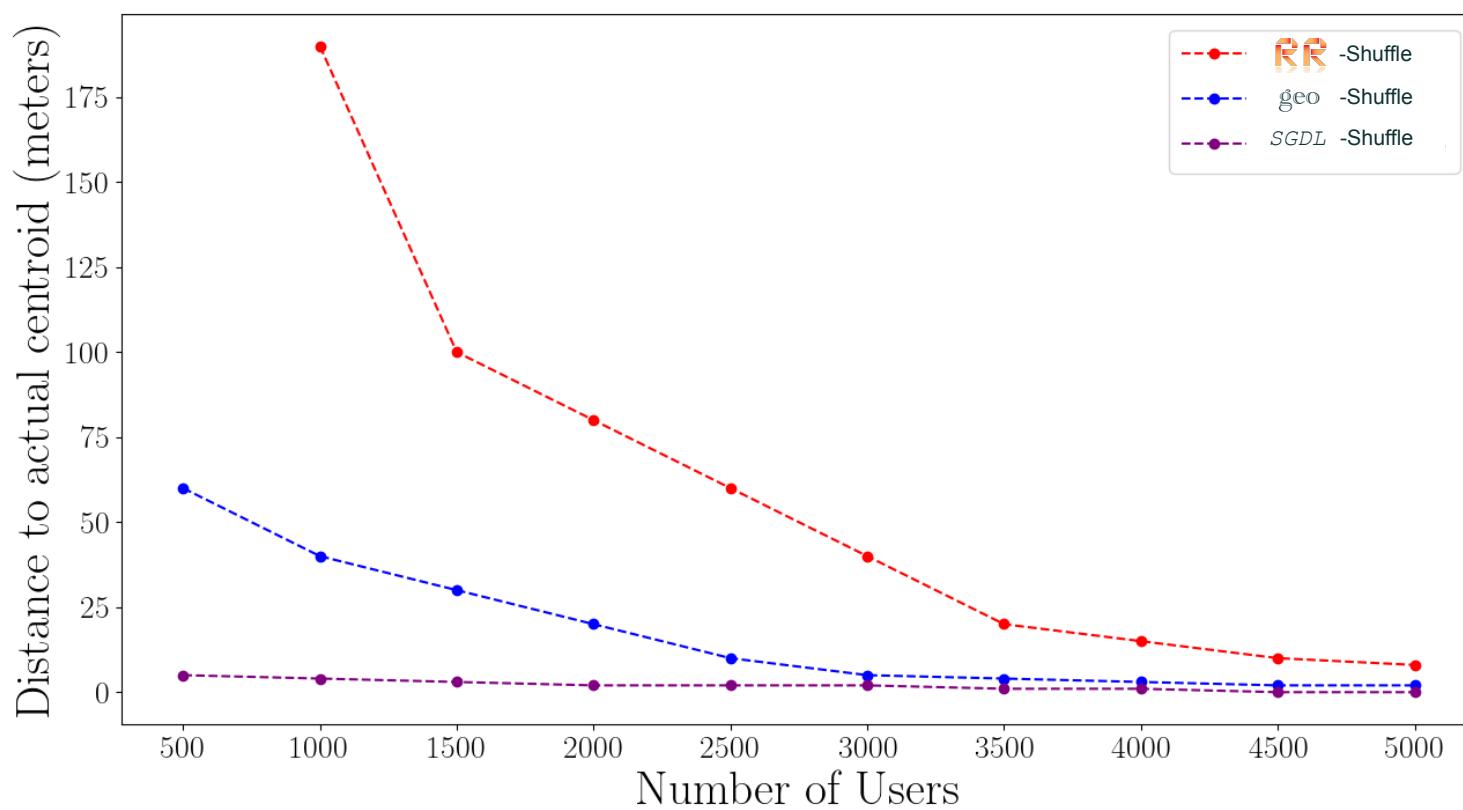
Utility Experiment

Find the centroid of addresses in Austin, Texas



$$\varepsilon = 0.15 \mid \delta = 10^{-4}$$

Utility Experiment: Results



Thank you :) Questions?

Andreas Athanasiou
andreas.athanasiou@inria.fr

